

# Global Bilateral Migration

*Remy LeWinter & Benjamin Romine*

## Setup

An installation of Neo4j is required to create this report.

Click the link below and follow the instructions to install Neo4j on your machine. <https://neo4j.com/docs/operations-manual/current/installation/> You should now be able to run Neo4j from the command line.

Now, download Neo4j Desktop from this address: <https://neo4j.com/download/> It will consist of only one file. Execute the file and Neo4j Desktop will open.

Create an account, log in, and create a project. Select the project and click “Add Graph”. Name the graph.

Click “Manage” and go to the Settings tab. Scroll to the line of code “dbms.directories.import=import”, it should be the first uncommented line. Comment it out by typing a “#” in front of it. Scroll to the line of code “dbms.security.auth\_enabled=true”, it should be the next uncommented line. Change it to “false” and click “Apply”.

Now click “Start” (the play button) on the graph and select “Continue Anyway”. Go to the Details tab and write down the HTTP port being used.

In this document, go to the first code chunk in the section labeled “Populating the Database” and ensure that the port number is correct (i.e. the number after “http://localhost:”).

Once the database is populated via the code in this document, you can open Neo4j Browser from the Desktop to directly visualise the data model and perform Cypher queries.

## Project Motivation

This project is designed as an exercise in using the property graph data model for highly connected data and analysing flow over a network. We use the Neo4j graph database to store and query bilateral migration and population data for over 200 countries.

## About the Data

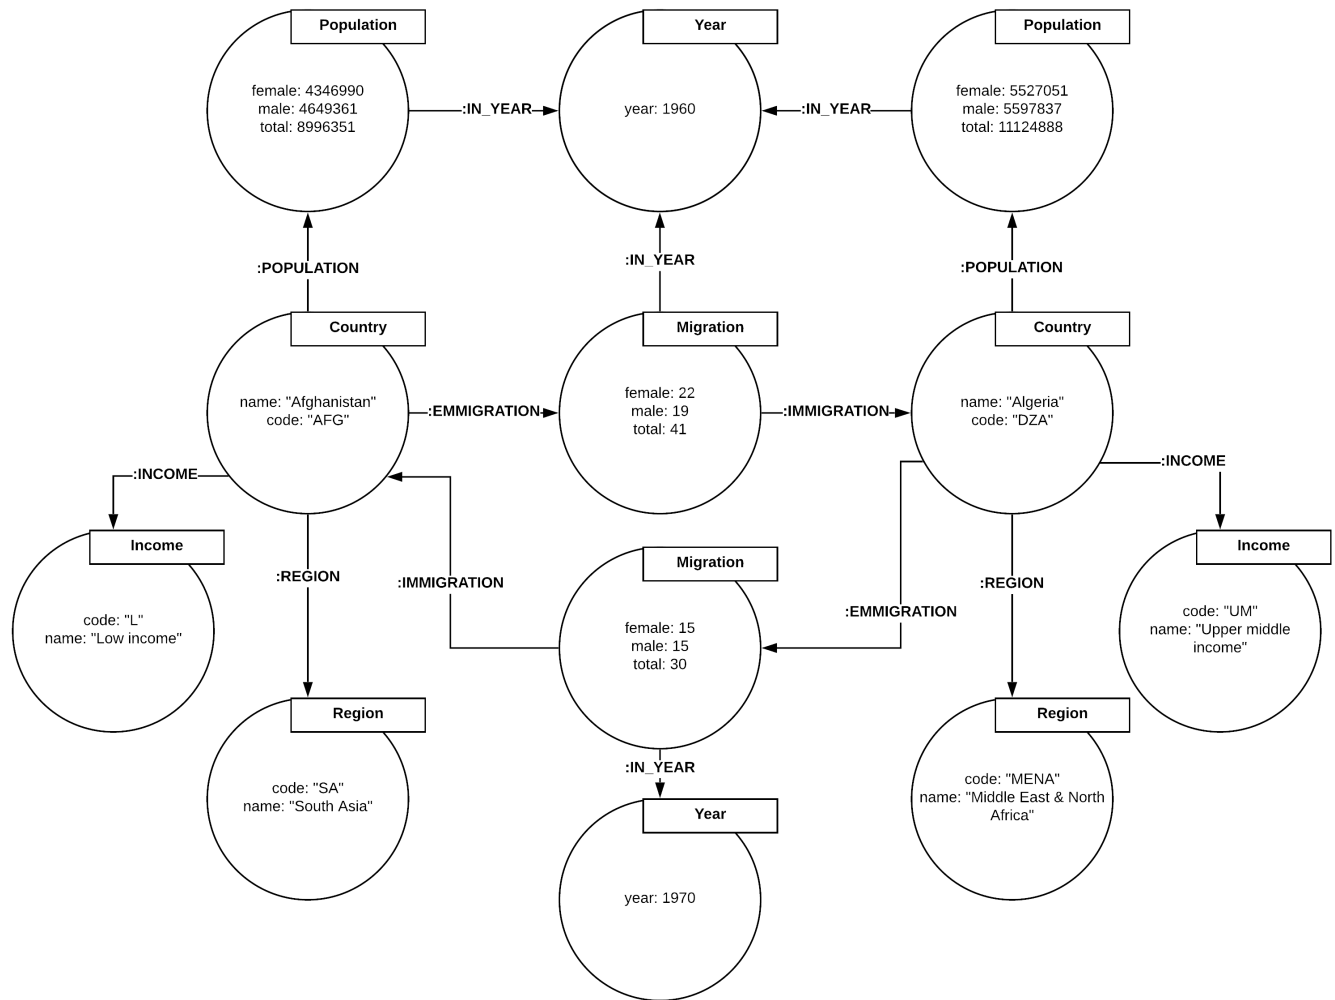
### Source:

This project uses all available data from the World Bank Global Bilateral Migration Database, containing the origin, destination, and gender of migrants for 232 countries in 1960, 1970, 1980, 1990, and 2000 with raw counts of annual migrant flow. This is joined with World Bank’s World Development Indicator country-level population data for those years. Each country is also associated with a region and a categorical income level in metadata supplied by World Bank.

### Data Model:

The property graph data model requires different thinking than the relational model. In the literal sense, relationships are more central to the graph model than the relational. Neo4j’s native graph storage directly encodes relationships between entities, as opposed to the relational model’s use of keys for joining tables.

Neo4j's native graph processing uses pointer chasing to match patterns of nodes and edges in relation some set of anchored nodes, which is much faster than searching on data contained in node and edge properties. Accordingly, facts are represented as nodes to optimise query performance. The chart below shows an example of our data model:



Data, metadata, and descriptions from World Bank are found in the `./data` folder. Processed data used in the database are in the `./data/final` folder.

## Processing Data and Populating the Database

After following the setup instructions above, connect to the database:

```
con <- neo4j_api$new(  
  url = "http://localhost:7474",  
  user = "neo4j",  
  password = "neo4j"  
)
```

Clear the database if there's anything in it:

```
call_neo4j("MATCH (n) DETACH DELETE n", con)
```

```
## No data returned.
```

```
## list()
```

```
## attr("class")
```

```
## [1] "neo" "neo" "list"
```

## Pre-Processing

Read the data and rename columns:

```
BilateralMigration <- select(read.csv("./data/Bilateral_Migration.csv", na.strings=c("", "NA")), -4)
Population <- read.csv("./data/Population.csv", na.strings=c("", "NA"))
IncomeRegion <- select(read.csv("./data/Country_Metadata.csv", na.strings=c("", "NA")), 1, 3, 4)

colnames(BilateralMigration) <- c("Origin_Name", "Origin_Code", "Migration_by_Gender", "Dest_Name", "Dest_Code")
colnames(Population) <- c("Series_Name", "Series_Code", "Country_Name", "Country_Code", "y1960", "y1970", "y1980", "y1990", "y2000")
colnames(IncomeRegion) <- c("Country_Code", "Income_Name", "Region_Name")
```

## Processing the Migration Data

If you run the commented line of code for each of the years, you will see that rows with missing values have no data. Therefore we select only rows with complete cases:

```
#BilateralMigration[!complete.cases(BilateralMigration),] %>% summarise(non_na = sum(!is.na(y2000)))

# select complete cases
BilateralMigration <- BilateralMigration[complete.cases(BilateralMigration),]

# gather the years and spread the genders
BilateralMigration <- gather(BilateralMigration, key="Year", value="Mig", y1960, y1970, y1980, y1990, y2000)
BilateralMigration$Year <- BilateralMigration$Year %>% str_sub(2, 5)
BilateralMigration <- spread(BilateralMigration, Migration_by_Gender, Mig)
```

## Processing the Population Data

Of the rows with missing values in Population, 6 have data for year 1990, and 7 have data for year 2000. Therefore we will not filter for only complete cases. Run the commented line below for each year to verify:

```
#Population[!complete.cases(Population),] %>% summarise(non_na = sum(!is.na(y2000)))
```

Replace “Series\_Name” and “Series\_Code” in Population with a single field:

```
Population <- Population %>%
  mutate(Population_by_Gender = Series_Code %>%
    str_replace("SP\\.POP\\.TOTL\\.MA\\.IN", "Male") %>%
    str_replace("SP\\.POP\\.TOTL\\.FE\\.IN", "Female") %>%
    str_replace("SP\\.POP\\.TOTL", "Total")) %>%
  select(Country_Name, Country_Code, Population_by_Gender, y1960, y1970, y1980, y1990, y2000)
```

Gather the years and spread the genders:

```

# gather
Population <- gather(Population, key="Year", value="Pop", y1960, y1970, y1980, y1990, y2000)

# remove prefixes
Population$Year <- Population$Year %>% str_sub(2, 5)

# we can take complete cases now that years are gathered
Population <- Population[complete.cases(Population),]

# spread
Population <- spread(Population, Population_by_Gender, Pop)

```

Filter Population to only contain data on countries found in BilateralMigration. It is ok if there is migration but not population data for some countries. Since all countries in BilateralMigration are both origins and destinations (run commented line to verify), we only need use one:

```

#setdiff(BilateralMigration$Dest_Code, BilateralMigration$Origin_Code)

Population <- Population %>% filter(Country_Code %in% BilateralMigration$Origin_Code)

```

## Processing Income and Region Data

IncomeRegion contains rows with missing data. All of these are for aggregates and demographic groups, so we'll take only complete cases:

```
IncomeRegion <- IncomeRegion[complete.cases(IncomeRegion),]
```

Create codes for income and region data:

```

IncomeRegion$Income_Code <- IncomeRegion$Income_Name %>%
  str_replace("High income", "H") %>%
  str_replace("Low income", "L") %>%
  str_replace("Lower middle income", "LM") %>%
  str_replace("Upper middle income", "UM")
IncomeRegion$Region_Code <- IncomeRegion$Region_Name %>%
  str_replace("East Asia & Pacific", "EAP") %>%
  str_replace("Europe & Central Asia", "ECA") %>%
  str_replace("Latin America & Caribbean", "LAC") %>%
  str_replace("Middle East & North Africa", "MENA") %>%
  str_replace("North America", "NA") %>%
  str_replace("South Asia", "SA") %>%
  str_replace("Sub-Saharan Africa", "SAA")

```

If you run the two commented lines below, you'll see that the same 8 rows in IncomeRegion do not appear in Population or BilateralMigration. Remove those rows:

```

#setdiff(IncomeRegion$Country_Code, BilateralMigration$Origin_Code)
#setdiff(IncomeRegion$Country_Code, Population$Country_Code)

IncomeRegion <- IncomeRegion %>% filter(Country_Code %in% Population$Country_Code)

```

## Save Modified Data to CSV

```
write.csv(BilateralMigration, file="./data/final/Bilateral_Migration_Final.csv", row.names=F)
write.csv(BilateralMigration %>% distinct(Origin_Name, Origin_Code), file="./data/final/Country_Final.csv", row.names=F)
write.csv(Population, file="./data/final/Population_Final.csv", row.names=F)
write.csv(IncomeRegion, file="./data/final/Income_Region_Final.csv", row.names=F)
```

## Populating Neo4j

Constraints:

```
# Some constraints are restricted to the Enterprise version and may not always work
"CREATE CONSTRAINT ON (c:Country) ASSERT exists(c.name)
CREATE CONSTRAINT ON (c:Country) ASSERT exists(c.code)
CREATE CONSTRAINT ON (c:Country) ASSERT c.name IS UNIQUE
CREATE CONSTRAINT ON (c:Country) ASSERT c.code IS UNIQUE

CREATE CONSTRAINT ON (m:Migration) ASSERT exists(m.total)
CREATE CONSTRAINT ON (m:Migration) ASSERT exists(m.male)
CREATE CONSTRAINT ON (m:Migration) ASSERT exists(m.female)

CREATE CONSTRAINT ON (y:Year) ASSERT exists(y.year)
CREATE CONSTRAINT ON (y:Year) ASSERT y.year IS UNIQUE

CREATE CONSTRAINT ON (i:Income) ASSERT exists(i.code)
CREATE CONSTRAINT ON (i:Income) ASSERT exists(i.name)
CREATE CONSTRAINT ON (i:Income) ASSERT i.code IS UNIQUE
CREATE CONSTRAINT ON (i:Income) ASSERT i.name IS UNIQUE

CREATE CONSTRAINT ON (r:Region) ASSERT exists(r.code)
CREATE CONSTRAINT ON (r:Region) ASSERT exists(r.name)
CREATE CONSTRAINT ON (r:Region) ASSERT r.code IS UNIQUE
CREATE CONSTRAINT ON (r:Region) ASSERT r.name IS UNIQUE;" %>%
  call_neo4j(con)
```

Years:

```
"CREATE (:Year {year: 1960}), (:Year {year: 1970}), (:Year {year: 1980}), (:Year {year: 1990}), (:Year {year: 2000})" %>%
  call_neo4j(con)
```

```
## No data returned.
```

```
## list()
## attr("class")
## [1] "neo" "neo" "list"
```

Countries:

```
on_load_country <- 'CREATE (:Country {code: csvLine.Origin_Code, name: csvLine.Origin_Name})'

country_path <- str_c("file://", getwd(), "/data/final/Country_Final.csv")

load_csv(url=country_path, con=con, on_load=on_load_country, as="csvLine", periodic_commit=500)
```

```
## No data returned.
```

```
## # A tibble: 12 x 2
##   type          value
##   <chr>        <dbl>
## 1 contains_updates      1
## 2 nodes_created      226
## 3 nodes_deleted        0
## 4 properties_set      452
## 5 relationships_created  0
## 6 relationship_deleted  0
## 7 labels_added      226
## 8 labels_removed        0
## 9 indexes_added        0
## 10 indexes_removed      0
## 11 constraints_added    0
## 12 constraints_removed  0
```

Migrations:

```
on_load_mig <-
'MATCH (o:Country {code: csvLine.Origin_Code})
MATCH (d:Country {code: csvLine.Dest_Code})
MATCH (y:Year {year: toInteger(csvLine.Year)})
CREATE (m:Migration {female: toInteger(csvLine.Female), male: toInteger(csvLine.Male), total: toInteger(csvLine.Total)})
MERGE (o)-[:EMMIGRATION]->(m)-[:IMMIGRATION]->(d)
MERGE (m)-[:IN_YEAR]->(y);'

mig_path <- str_c("file://", getwd(), "/data/final/Bilateral_Migration_Final.csv")

load_csv(url=mig_path, con=con, on_load=on_load_mig, as="csvLine", periodic_commit=500)
```

## No data returned.

```
## # A tibble: 12 x 2
##   type          value
##   <chr>        <dbl>
## 1 contains_updates      1
## 2 nodes_created    255380
## 3 nodes_deleted        0
## 4 properties_set    766140
## 5 relationships_created 766140
## 6 relationship_deleted  0
## 7 labels_added    255380
## 8 labels_removed        0
## 9 indexes_added        0
## 10 indexes_removed      0
## 11 constraints_added    0
## 12 constraints_removed  0
```

Populations:

```
on_load_pop <-
'MATCH (c:Country {code: csvLine.Country_Code})
MATCH (y:Year {year: toInteger(csvLine.Year)})
MERGE (p:Population {total: toInteger(csvLine.Total)})
FOREACH(ignoreMe IN CASE WHEN trim(csvLine.Female) <> "" THEN [1] ELSE [] END | SET p.female = toInteger(csvLine.Female))
FOREACH(ignoreMe IN CASE WHEN trim(csvLine.Male) <> "" THEN [1] ELSE [] END | SET p.male = toInteger(csvLine.Male))
MERGE (c)-[:POPULATION]->(p)-[:IN_YEAR]->(y);'
```

```
pop_path <- str_c("file://", getwd(), "/data/final/Population_Final.csv")

load_csv(url=pop_path, con=con, on_load=on_load_pop, as="csvLine", periodic_commit=500)
```

```
## No data returned.
```

```
## # A tibble: 12 x 2
##   type          value
##   <chr>         <dbl>
## 1 contains_updates      1
## 2 nodes_created      1036
## 3 nodes_deleted        0
## 4 properties_set      2920
## 5 relationships_created 2074
## 6 relationship_deleted   0
## 7 labels_added       1036
## 8 labels_removed        0
## 9 indexes_added        0
## 10 indexes_removed      0
## 11 constraints_added     0
## 12 constraints_removed   0
```

Incomes and Regions:

```
on_load_ir <-
'MATCH (c:Country {code: csvLine.Country_Code})
MERGE (i:Income {code: csvLine.Income_Code, name: csvLine.Income_Name})
MERGE (r:Region {code: csvLine.Region_Code, name: csvLine.Region_Name})
MERGE (c)-[:INCOME]->(i)
MERGE (c)-[:REGION]->(r);'

ir_path <- str_c("file://", getwd(), "/data/final/Income_Region_Final.csv")

load_csv(url=ir_path, con=con, on_load=on_load_ir, as="csvLine", periodic_commit=500)
```

```
## No data returned.
```

```
## # A tibble: 12 x 2
##   type          value
##   <chr>         <dbl>
## 1 contains_updates      1
## 2 nodes_created        11
## 3 nodes_deleted        0
## 4 properties_set       22
## 5 relationships_created 416
## 6 relationship_deleted   0
## 7 labels_added         11
## 8 labels_removed        0
## 9 indexes_added        0
## 10 indexes_removed      0
## 11 constraints_added     0
## 12 constraints_removed   0
```

Now that the database is populated, let's free up some RAM:

```
rm(BilateralMigration, IncomeRegion, Population)
```

## Summary Statistics

### Number of Countries in each Income Level

```
countries_per_income_level <- data.frame(call_neo4j("match (i:Income)<-[ic:INCOME]-(:Country) return i.",
colnames(countries_per_income_level) <- c("Income Level", "Number of Countries")
print(countries_per_income_level)
```

##	Income Level	Number of Countries
## 1	Low income	33
## 2	Upper middle income	54
## 3	High income	75
## 4	Lower middle income	46

### Number of Countries in each Region

```
countries_per_region <- data.frame(call_neo4j("match (r:Region)<-[rc:REGION]-(:Country) return r.name, c",
colnames(countries_per_region) <- c("Region", "Number of Countries")
print(countries_per_region)
```


##	Region	Number of Countries
## 1	South Asia	8
## 2	Europe & Central Asia	53
## 3	Middle East & North Africa	21
## 4	East Asia & Pacific	37
## 5	Sub-Saharan Africa	47
## 6	Latin America & Caribbean	39
## 7	North America	3

### Yearly Global Population

```
global_pop <- call_neo4j("match (p:Population)-->(y:Year) return y.year as Year, sum(p.female) as Female, sum(p.male) as Male")
colnames(global_pop) <- c("Year", "Female", "Male")
global_pop <- gather(global_pop, `Female`, `Male`, key="Gender", value="Population")

ggplot(global_pop, aes(x=`Year`, y=`Population`, fill=`Gender`)) +
  geom_bar(position="stack", stat="identity")
```





### Yearly Global Migration

```
global_mig <- call_neo4j("match (m:Migration)-->(y:Year) return y.year as Year, sum(m.female) as Female")
colnames(global_mig) <- c("Year", "Female", "Male")
global_mig <- gather(global_mig, `Female`, `Male`, key="Gender", value="Migrants")

ggplot(global_mig, aes(x=`Year`, y=`Migrants`, fill=`Gender`)) +
  geom_bar(position="stack", stat="identity")
```

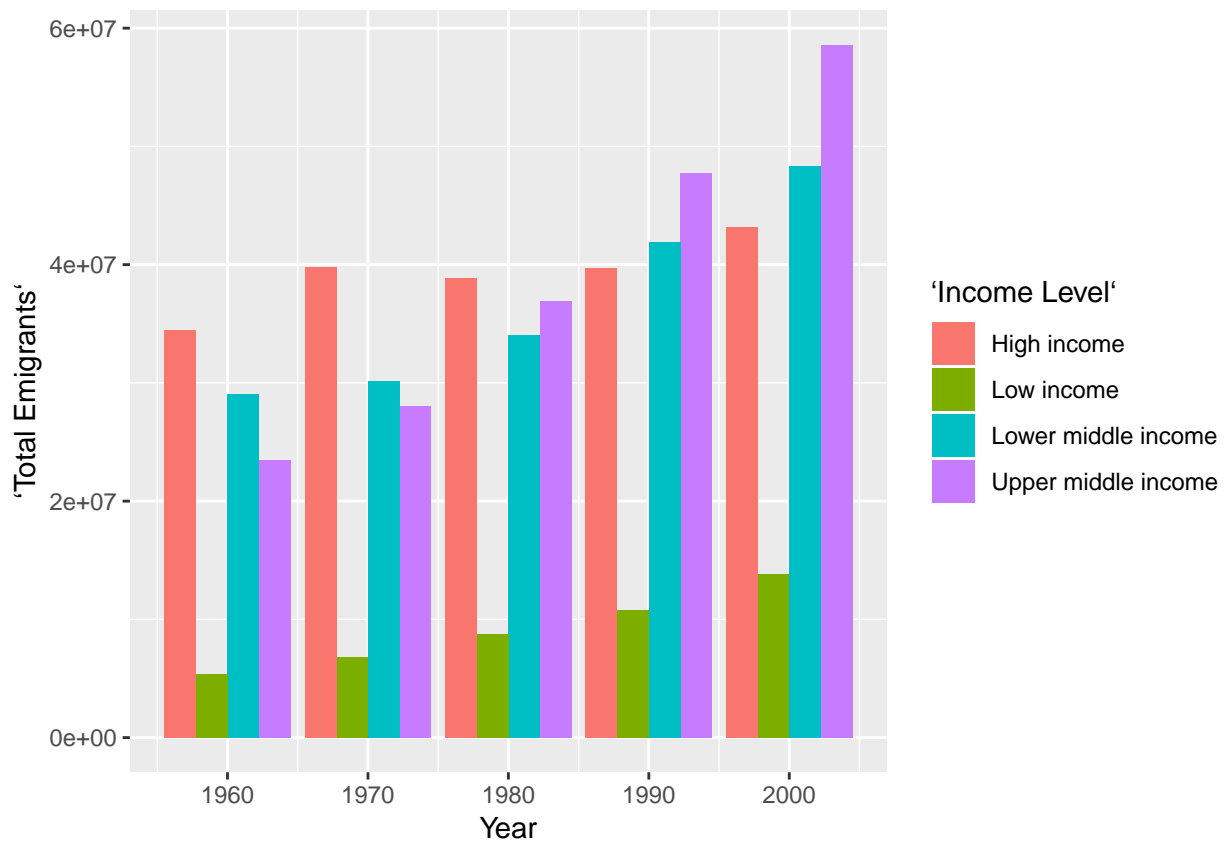


### Total Migration by Income Level and Gender

```
emigrants_income <- call_neo4j('MATCH (i:Income)<-[:INCOME]-(c:Country)-[:EMMIGRATION]->(m:Migration)-[
  data.frame()
colnames(emigrants_income) <- c("Income Level", "Year", "Total Emigrants", "Total Male Emigrants", "Total
immigrants_income <- call_neo4j('MATCH (i:Income)<-[:INCOME]-(c:Country)<-[:IMMIGRATION]-(m:Migration)-[
  data.frame()
colnames(immigrants_income) <- c("Income Level", "Year", "Total Immigrants", "Total Male Immigrants", "Total
```

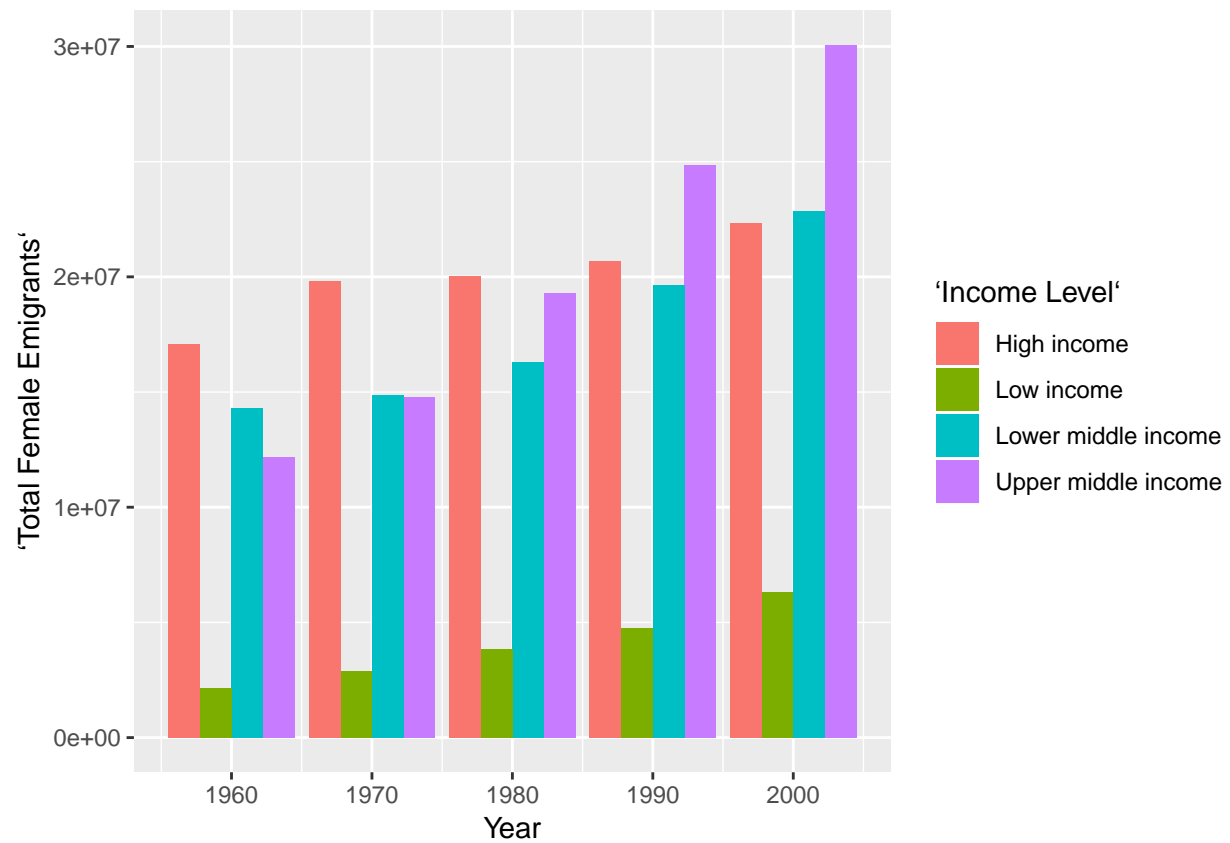
Over time, emigration has increased for all income levels, but the increase is most drastic for middle income countries:

```
ggplot(emigrants_income, aes(x = `Year`, y = `Total Emigrants`, fill = `Income Level`)) +
  geom_bar(position = "dodge", stat = "identity")
```




Emigration was roughly proportional between women and men, though the proportion of men emigrating from upper-middle income countries took longer to surpass emigration from lower-middle income countries than it did for women, for whom the difference is much greater:

```
ggplot(emigrants_income, aes(x = `Year`, y = `Total Female Emigrants`, fill = `Income Level`)) +
  geom_bar(position = "dodge", stat = "identity")
```




```
ggplot(emigrants_income, aes(x = `Year`, y = `Total Male Emigrants`, fill = `Income Level`)) +
  geom_bar(position = "dodge", stat = "identity")
```




High income countries always received the lion's share of immigrants, and they seem to be the destination for most of the increase in migrants shown in the previous charts, such that their share grew larger every decade. Immigration to upper-middle income countries experienced a much less dramatic rise each decade, while immigration to lower-middle income countries actually decreased between 1960 and 2000:

```
ggplot(immigrants_income, aes(x = `Year`, y = `Total Immigrants`, fill = `Income Level`)) +  
  geom_bar(position = "dodge", stat = "identity")
```




Immigration to different income-level countries followed roughly the same trends for men and women:

```
ggplot(immigrants_income, aes(x = `Year`, y = `Total Female Immigrants`, fill = `Income Level`)) +  
  geom_bar(position = "dodge", stat = "identity")
```



```
ggplot(immigrants_income, aes(x = `Year`, y = `Total Male Immigrants`, fill = `Income Level`)) +
  geom_bar(position = "dodge", stat = "identity")
```



Now, let's try using this data to predict the total number of immigrants to all countries with different income levels in 2010. First, we create linear regression models:

```
immigrant_income_linreg <- lm(`Total Immigrants` ~ `Income Level` + Year, immigrants_income)
emigrant_income_linreg <- lm(`Total Emigrants` ~ `Income Level` + Year, emigrants_income)
income_levels <- c("Low income", "Lower middle income", "Upper middle income", "High income")
new_data_income <- data.frame("Income Level" = income_levels, "Year" = rep(2010, 4))
colnames(new_data_income) <- c("Income Level", "Year")
summary(immigrant_income_linreg)
```

```
##
## Call:
## lm(formula = `Total Immigrants` ~ `Income Level` + Year, data = immigrants_income)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-18433135	-4810892	-632322	4739676	27835865

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-830900433	353854144	-2.348	0.0330
`Income Level`Low income	-62628413	7147839	-8.762	2.76e-07
`Income Level`Lower middle income	-39954819	7147839	-5.590	5.16e-05
`Income Level`Upper middle income	-46070873	7147839	-6.445	1.10e-05
Year	454220	178696	2.542	0.0226

```
##
## (Intercept) *
```

`Income Level`Low income	***



```
## `Income Level`Lower middle income ***
## `Income Level`Upper middle income ***
## Year *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 11300000 on 15 degrees of freedom
## Multiple R-squared:  0.8562, Adjusted R-squared:  0.8179
## F-statistic: 22.33 on 4 and 15 DF,  p-value: 3.576e-06
```

```
summary(emigrant_income_linreg)
```

```
##
## Call:
## lm(formula = `Total Emigrants` ~ `Income Level` + Year, data = emigrants_income)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -6579114 -3094973 -313376  3081755 10702913
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      -843952398   158259952  -5.333 8.37e-05
## `Income Level`Low income      -30087071    3196845  -9.411 1.10e-07
## `Income Level`Lower middle income  -2478609    3196845  -0.775  0.450
## `Income Level`Upper middle income  -235356    3196845  -0.074  0.942
## Year              446033      79921    5.581 5.25e-05
##
## (Intercept) ***
## `Income Level`Low income ***
## `Income Level`Lower middle income
## `Income Level`Upper middle income
## Year ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5055000 on 15 degrees of freedom
## Multiple R-squared:  0.9127, Adjusted R-squared:  0.8895
## F-statistic: 39.22 on 4 and 15 DF,  p-value: 8.938e-08
```

Notice that the residual standard error values are very high. Using these linear regression models, we will create confidence intervals for the values that we will predict:

```
predictions_2010_income <- data.frame(income_levels, predict(immigrant_income_linreg, new_data_income,
colnames(predictions_2010_income) <-
  c("Income Level", "Expected Immigrants in 2010", "Lower Bound", "Upper Bound", "Expected Emigrants in
predictions_2010_income
```

```
##      Income Level Expected Immigrants in 2010 Lower Bound Upper Bound
## 1      Low income      19453183      3749038      35157329
## 2 Lower middle income      42126777      26422632      57830922
## 3 Upper middle income      36010723      20306578      51714868
## 4      High income      82081596      66377451      97785742
## Expected Emigrants in 2010 Lower Bound Upper Bound
## 1      22486902      15463281      29510522
## 2      50095363      43071743      57118984
```

```
## 3          52338616    45314996    59362237
## 4          52573973    45550352    59597593
```

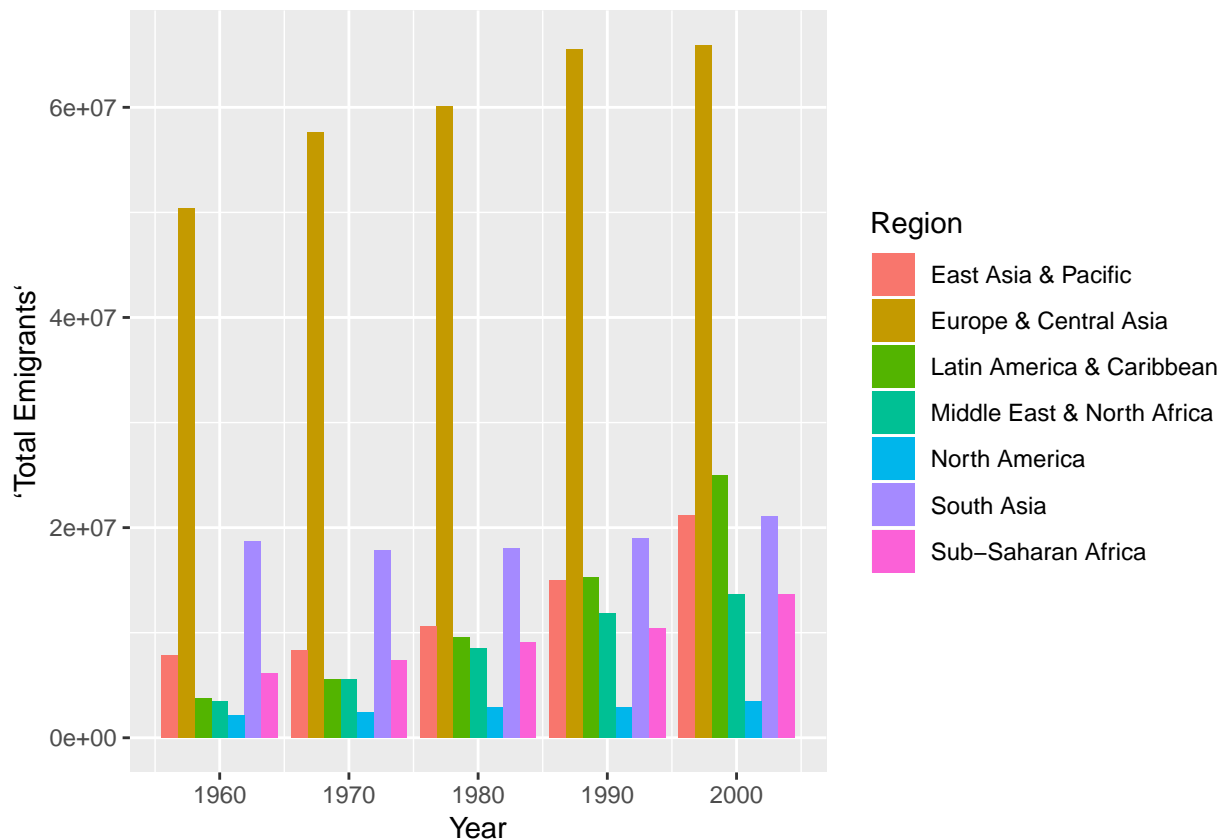
The confidence intervals for immigration are very large, so the data that is presented in these linear regressions is not much better than wild guessing.

### Total Migration by Region and Gender

```
emigrants_region <- call_neo4j('MATCH (r:Region)-[:REGION]-(c:Country)-[:EMMIGRATION]->(m:Migration)-[
  data.frame()
colnames(emigrants_region) <- c("Region", "Year", "Total Emigrants", "Total Male Emigrants", "Total Female Emigrants")
immigrants_region <- call_neo4j('MATCH (r:Region)-[:REGION]-(c:Country)-[:IMMIGRATION]->(m:Migration)-[
  data.frame()
colnames(immigrants_region) <- c("Region", "Year", "Total Immigrants", "Total Male Immigrants", "Total Female Immigrants")
```


Emigrants leave Europe and Central Asia at a much greater rate than any other region. While emigration from South Asia and North America stayed roughly equal over time, emigration from every other region increased each decade.

```
ggplot(emigrants_region, aes(x = `Year`, y = `Total Emigrants`, fill = `Region`)) +
  geom_bar(position = "dodge", stat = "identity")
```

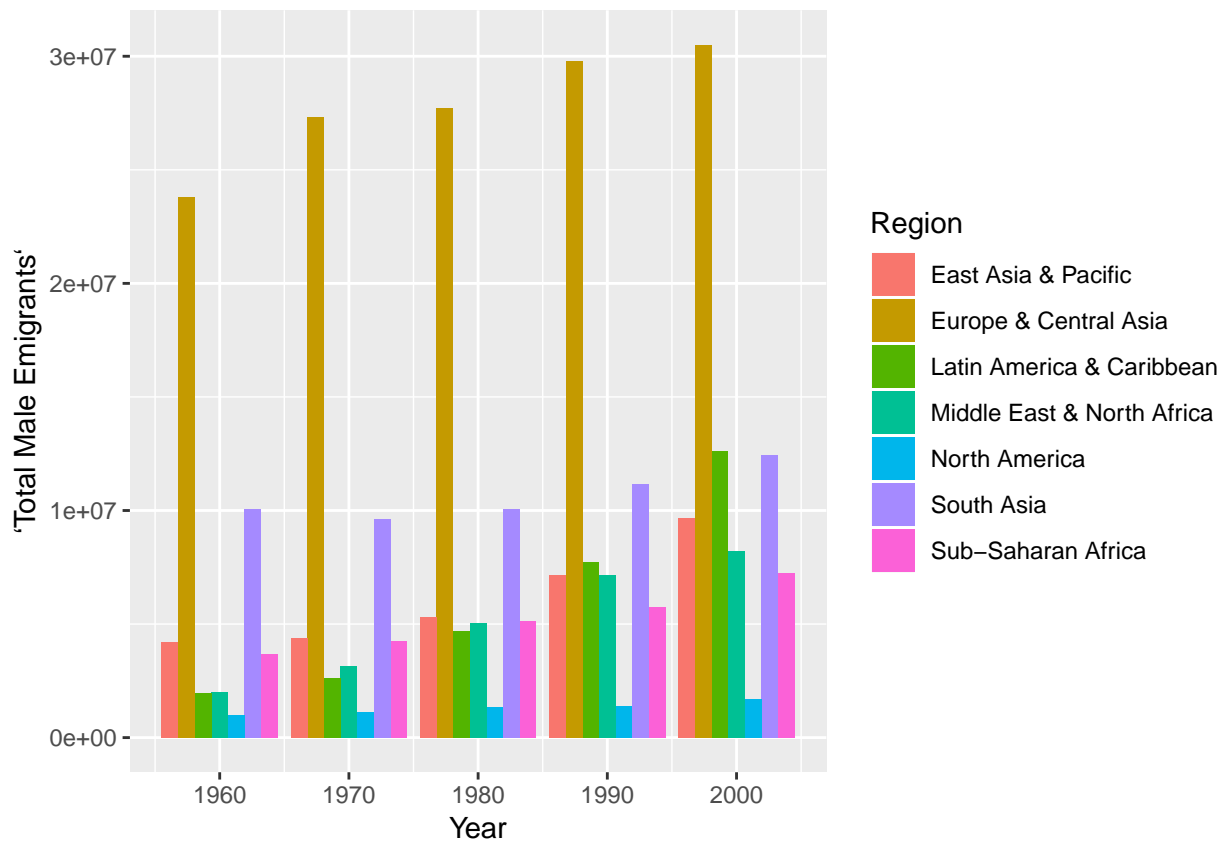


Women and men seem to emigrate from the different regions at proportional rates, with one exception - South Asian emigrants are more likely to be male:

```
ggplot(emigrants_region, aes(x = `Year`, y = `Total Female Emigrants`, fill = `Region`)) +
  geom_bar(position = "dodge", stat = "identity")
```




```
ggplot(emigrants_region, aes(x = `Year`, y = `Total Male Emigrants`, fill = `Region`)) +
  geom_bar(position = "dodge", stat = "identity")
```




Immigration to Latin America and Caribbean countries and Sub-Saharan Africa stayed roughly equal over time. Immigration to South Asia actually decreased every decade, while immigration to Europe and Central Asia increased monotonically and the proportion immigrating to North America greatly increased from 1980-2000:

```
ggplot(immigrants_region, aes(x = `Year`, y = `Total Immigrants`, fill = `Region`)) +
  geom_bar(position = "dodge", stat = "identity")
```




Women and men seem to immigrate to the different regions at proportional rates, with one exception - men are much more likely to immigrate to the Middle East and North Africa:

```
ggplot(immigrants_region, aes(x = `Year`, y = `Total Female Immigrants`, fill = `Region`)) +  
  geom_bar(position = "dodge", stat = "identity")
```



```
ggplot(immigrants_region, aes(x = `Year`, y = `Total Male Immigrants`, fill = `Region`)) +
  geom_bar(position = "dodge", stat = "identity")
```



As before, we're going to try to predict the total number of migrants to and from each region of the world for the year 2010. We create linear regressions once more:

```
immigrant_region_linreg <- lm(`Total Immigrants` ~ `Region` + Year, immigrants_region)
emigrant_region_linreg <- lm(`Total Emigrants` ~ `Region` + Year, emigrants_region)
regions <- unique(emigrants_region$Region)
new_data_region <- data.frame("Region" = regions, "Year" = rep(2010, 7))
colnames(new_data_income) <- c("Region", "Year")
summary(immigrant_region_linreg)
```

```
##
## Call:
## lm(formula = `Total Immigrants` ~ Region + Year, data = immigrants_region)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -11579234 -2977989   -38173   2553975  11759636
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -503003593  127765596  -3.937  0.000523
## RegionEurope & Central Asia    40980575    3413896  12.004  2.46e-12
## RegionLatin America & Caribbean  -5012983    3413896  -1.468  0.153549
## RegionMiddle East & North Africa -1440839    3413896  -0.422  0.676329
## RegionNorth America    12522218    3413896   3.668  0.001058
## RegionSouth Asia         3576277    3413896   1.048  0.304127
## RegionSub-Saharan Africa  -1855935    3413896  -0.544  0.591149
## Year              259554         64517   4.023  0.000416
```

```
##
## (Intercept) ***
## RegionEurope & Central Asia ***
## RegionLatin America & Caribbean
## RegionMiddle East & North Africa
## RegionNorth America **
## RegionSouth Asia
## RegionSub-Saharan Africa
## Year ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5398000 on 27 degrees of freedom
## Multiple R-squared:  0.9122, Adjusted R-squared:  0.8894
## F-statistic: 40.06 on 7 and 27 DF,  p-value: 1.207e-12
```

```
summary(emigrant_income_linreg)
```

```
##
## Call:
## lm(formula = `Total Emigrants` ~ `Income Level` + Year, data = emigrants_income)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -6579114 -3094973 -313376  3081755 10702913
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      -843952398   158259952  -5.333 8.37e-05
## `Income Level`Low income      -30087071    3196845  -9.411 1.10e-07
## `Income Level`Lower middle income  -2478609    3196845  -0.775  0.450
## `Income Level`Upper middle income  -235356    3196845  -0.074  0.942
## Year              446033      79921    5.581 5.25e-05
##
## (Intercept) ***
## `Income Level`Low income ***
## `Income Level`Lower middle income
## `Income Level`Upper middle income
## Year ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5055000 on 15 degrees of freedom
## Multiple R-squared:  0.9127, Adjusted R-squared:  0.8895
## F-statistic: 39.22 on 4 and 15 DF,  p-value: 8.938e-08
```

Similarly, the residual standard errors are very large. However, we will still attempt to predict the numbers of migrants for the year 2010:

```
predictions_2010_region <- data.frame(regions, predict(immigrant_region_linreg, new_data_region, interval="prediction"),
colnames(predictions_2010_region) <-
  c("Region", "Expected Immigrants in 2010", "Lower Bound", "Upper Bound", "Expected Emigrants in 2010")
predictions_2010_region
```

```
##              Region Expected Immigrants in 2010 Lower Bound
## 1              South Asia              22276701              15928121
```



## 2	Europe & Central Asia	59680999	53332418
## 3	Middle East & North Africa	17259585	10911005
## 4	East Asia & Pacific	18700424	12351844
## 5	Sub-Saharan Africa	16844489	10495908
## 6	Latin America & Caribbean	13687441	7338861
## 7	North America	31222642	24874062
##	Upper Bound Expected Emigrants in 2010	Lower Bound	Upper Bound
## 1	28625281	26575053	23057184
## 2	66029579	67548824	64030955
## 3	23608165	16273435	12755566
## 4	25049004	20241628	16723758
## 5	23193069	16988058	13470189
## 6	20036021	19458509	15940640
## 7	37571222	10409348	6891478
			13927217

Once again, it is clear that the confidence intervals for these predicted values are very large, and the linear regression models that we created are not very accurate.

## Limitations & Challenges

Below are listed some of the limitations that we discovered when it came to performing analysis on this data:

- The data only goes up to the year 2000, and does not account for significant historical events in recent history that have had a deal of impact on migration patterns (9/11, travel bans, international politics).
- The data is in 10-year increments, and therefore year-to-year prediction was not possible.
- When attempting to create linear regressions for the data, there were very large residual standard error values, which resulted in inaccurate predictions with little confidence.

## Future Work

One of the possible ways to utilize this data in the future is to link it to a database of news articles that go into detail about international politics to show the different ways that certain events affected immigration patterns in the future. Our chosen graph model is designed to make best use of the pointer-chasing employed by neo4j's native graph storage and processing. It is worth noting that our strategy of modeling all facts as nodes, i.e.  $(:Country)-[:MIGRATION]->(:Country)$  actually becomes a hindrance when using graph algorithms for things like clustering, centrality, link prediction, etc., because they must treat nodes of interest (e.g. countries) as directly adjacent.